

RAG-Augmented Multi-Agent LLM Framework for Explainable Software Vulnerability Detection Using CWE Knowledge Bases

Richa Verma (25MCSS02) Advisor: Dr. Akshay Pandey

How to run it, and what is still open

This report adds a local screen and a list of unfinished tests. It continues report 9.

What we did


A web page on this computer lets us pick a test set, a setting, and a token budget, then start a run.

The page only listens on this machine. It has no login because it is not open to the network.

The six public test sets are now stored in the project, with their licenses noted. A new copy of the project does not have to download them again.

GitHub runs the unit tests on each push. Those tests do not call the live model.

Picture



Numbers for this step only

Still open	Status
Finish the one Juliet file cut off by the token limit	Not finished
Juliet again with CWE search turned off	Not run
Sliced real bugs from Vul4J	Not run for the live model
A second model, if tokens remain	Not run
A person marking whether explanations are good	Labels still empty

What is next

The next real run needs a Groq key and the daily token budget. Empty human labels stay empty until a person fills them.

What this report does not claim

We still do not claim that CWE search improves Juliet detection. We have no score for agreement between human raters.